

# 16

## TCP/UDP Load Balancing with NGINX: Overview, Tips, and Tricks

Konstantin Pavlov, NGINX Inc.



# TCP load balancing

```
stream {  
    upstream mysql_backends {  
        server backend1.example.com:3306;  
        server backend2.example.com:3306;  
    }  
    server {  
        listen 3306;  
        proxy_pass mysql_backends;  
    }  
}
```



# UDP load balancing

```
stream {  
    upstream named_backends {  
        server named1.example.com:53;  
        server named2.example.com:53;  
    }  
    server {  
        listen 53 udp;  
        proxy_pass named_backends;  
        proxy_responses 1;  
        error_log logs/dns.log;  
    }  
}
```



# TCP/UDP load balancer tuning

## Supported load-balancing methods:

- weighted round-robin
- hashed key
- least number of connections

```
upstream backends-imaps {  
    hash $remote_addr consistent;  
    server 192.168.1.1:993;  
    server 192.168.1.2:993;  
}
```

## Additionally available in NGINX Plus:

- least time to connect
- least time to receive first byte
- least time to receive last byte

```
upstream backends-smtp {  
    least_time connect;  
    server 192.168.1.1:25;  
    server 192.168.1.2:25;  
}
```

# TCP/UDP load balancer tuning #2

```
upstream backends-ldap {
    server 192.168.1.1:389 weight=3 max_fails=2 fail_timeout=30;
    server 192.168.1.2:389 weight=1 max_fails=2 fail_timeout=30;
    server 192.168.1.3:389 down;
    server 192.168.1.4:389 backup;
}

upstream backends-pop3 {
    zone pop3-dynamic 64k;
    server 192.168.2.1:110 max_conns=20 slow_start=60s;
    server 192.168.2.2:110 max_conns=20 slow_start=60s;
}

upstream backends-webservers {
    zone web-dynamic 128k;
    server microservice1.example.com service=http resolve;
}
```



# TCP/UDP active health checks

```
stream {
    upstream backend-imaps {
        zone dynamic 64k;
        server 192.168.1.1:993;
    }
    server {
        listen 993;
        proxy_pass backend-imaps;
        health_check interval=10 passes=2 fails=3 port=8080 match=http;
        health_check_timeout 5s;
    }
    match http {
        send "GET /status HTTP/1.0\r\nHost: imaps.example.com\r\n\r\n";
        expect ~* "200 OK";
    }
}
```

# TCP/UDP active health checks #2

```
upstream dns-backends {
    zone dynamic 64k;
    server 192.168.1.100:53;
}

server {
    listen 53;
    listen 53 udp;
    proxy_pass dns-backends;
    health_check match=dns_test udp;
}

match dns_test {
    send
"\xaa\x67\x01\x00\x00\x01\x00\x00\x00\x00\x00\x05\x6e\x67\x69\x6e\x78\x03\x6f\x72\x6
7\x00\x00\x01\x00\x01"; # A? nginx.org
    expect ~ "\xce\xfb\xff\x3f"; # 206.251.255.63
}
```



# Access control and limiting

```
stream {  
    limit_conn_zone $binary_remote_addr zone=addr:10m;  
    server {  
        listen 5432;  
        deny 192.168.1.2;  
        allow 192.168.1.1/24;  
        allow 2001:0db8::/32;  
        deny all;  
        limit_conn addr 1;  
        proxy_download_rate 100k;  
        proxy_upload_rate 50k;  
        proxy_pass postgres.int.example.com:5432;  
    }  
}
```





# Passing client's IP to the backends

```
stream {  
    server {  
        listen 192.168.1.100:80;  
        proxy_protocol on;  
        proxy_pass 192.168.2.1:80;  
    }  
}
```



# Passing client's IP to the backends #2

```
user root;
```

```
# requires additional OS-level configuration!
```

```
# https://www.kernel.org/doc/Documentation/networking/tproxy.txt
```

```
stream {
```

```
    server {
```

```
        listen 192.168.1.1:1234;
```

```
        proxy_bind $remote_addr transparent;
```

```
        proxy_pass 192.168.2.1:1234;
```

```
    }
```

```
}
```



# TLS: termination

```
stream {  
    server {  
        listen 192.168.1.100:443 ssl;  
        ssl_certificate /var/lib/acme/live/thre.sh/fullchain;  
        ssl_certificate_key /var/lib/acme/live/thre.sh/privkey;  
        proxy_ssl off;  
        proxy_pass 192.168.1.101:80;  
    }  
}
```



# TLS: re-encryption

```
stream {  
    server {  
        listen 192.168.1.100:443 ssl;  
        ssl_certificate /var/lib/acme/live/thre.sh/fullchain;  
        ssl_certificate_key /var/lib/acme/live/thre.sh/privkey;  
        proxy_ssl on;  
        proxy_ssl_verify on;  
        proxy_ssl_trusted_certificate /etc/ssl/trusted.pem;  
        proxy_pass 192.168.1.101:443;  
    }  
}
```



# TLS: wrapping

```
stream {  
    server {  
        listen 192.168.1.100:80;  
        proxy_ssl on;  
        proxy_ssl_verify on;  
        proxy_ssl_trusted_certificate /etc/ssl/trusted.pem;  
        proxy_pass 192.168.1.101:443;  
    }  
}
```



# TCP/UDP load balancer: logging

```
2016/08/21 18:39:04 [info] 82529#100542: *4507759 client 192.168.55.64:2372 connected to 192.168.55.64:15431
```

```
2016/08/21 18:39:04 [info] 82529#100542: *4507759 proxy 192.168.55.64:2373 connected to 192.168.55.64:15434
```

```
2016/08/21 18:39:04 [info] 82529#100542: *4507759 upstream disconnected, bytes from/to client:106/7728, bytes from/to upstream:7728/106
```

```
2016/08/21 18:39:04 [info] 82529#100542: *4507752 udp client 192.168.55.64:46255 connected to 192.168.55.64:53
```

```
2016/08/21 18:39:04 [info] 82529#100542: *4507752 udp proxy 192.168.55.64:46256 connected to 192.168.55.64:53
```

```
2016/08/21 18:39:04 [info] 82529#100542: *4507754 udp upstream disconnected, bytes from/to client:27/126, bytes from/to upstream:126/27
```



# TCP/UDP load balancer: better logging

```
log_format combined '$remote_addr - - [$time_local] $protocol  
$status $bytes_sent $bytes_received $session_time "$upstream_addr"';
```

```
access_log /var/log/nginx/stream-access.log combined;
```

```
127.0.0.1 - - [26/Aug/2016:18:04:16 +0300] TCP 200 158 90 0.007  
"127.0.0.1:9000"
```



# TCP/UDP load balancer: variables!

Since NGINX 1.11.2:

- map to build variables using other variables
- geo module to build variables based on client's IP addresses
- geoip module to create variables using MaxMind GeoIP
- split\_clients to enable A/B testing
- nginxScript via js\_set



# TCP/UDP load balancer: variables!

```
stream {  
    server {  
        listen 2007;  
        listen [::]:2007 udp;  
        return $remote_addr;  
    }  
}
```

```
thresh@fruity ~ $ echo "foo" | nc -w 1 127.0.0.1 2007  
127.0.0.1  
thresh@fruity ~ $ echo "foo" | nc -6 -w 1 -u ::1 2007  
::1
```

# TCP/UDP load balancer: variables!

```
stream {  
    geoip_country GeoIP.dat;  
    geoip_city GeoLiteCity.dat;  
    limit_conn_zone $geoip_country_code zone=addr:10m;  
    split_clients "${remote_addr}AAA" $upstream {  
        0.5% 192.168.1.101:5000; # feature test  
        *    192.168.1.100:5000; # production  
    }  
    server {  
        proxy_pass $upstream;  
        limit_conn addr 5;  
    }  
}
```

# Extending TCP/UDP LB with nginxScript

nginx.conf:

```
load_module modules/nginx_stream_js_module.so;
stream {
    js_include stream.js;
    server {
        listen 8000;
        js_set $foo foo;
        return $foo;
    }
}
```

stream.js:

```
function foo(s) {
    s.log("hello from foo() handler!");
    return s.remoteAddress;
}
```



# Extending TCP/UDP LB with nginxScript

Moving the requests gradually between the backends: enjoy the demo!

nginx.conf-transition and transition.js:  
<https://github.com/thresheek/nginxconf16/>



# TCP/UDP payload fiddling w/nginxScript

Simple WAF to protect against ShellShock  
and HTTPoxy: enjoy the demo!

nginx.conf and waf.js:

<https://github.com/thresheek/nginxconf16/>



# TCP/UDP w/nginxScript: performance?

TODO: Some performance metrics



# TCP/UDP load balancer future

As the stream module is rather new, we welcome the feedback and feature requests you think would be valuable for your use cases.



## See also

TCP/UDP Load balancer admin guide:

<https://www.nginx.com/resources/admin-guide/tcp-load-balancing/>

Load Balancing TCP traffic:

<https://www.nginx.com/blog/advanced-mysql-load-balancing-with-nginx-plus/>

Load Balancing DNS traffic:

<https://www.nginx.com/blog/load-balancing-dns-traffic-nginx-plus/>

nginxScript documentation:

<http://hg.nginx.org/njs/file/tip/README>





# Thank you!

Questions?

Configuration snippets and slides:

<http://github.com/thresheek/somewhere/>

Konstantin Pavlov, Nginx Inc.

@cryothresh